

Red-Black Trees-I

ESO207

Indian Institute of Technology, Kanpur

Motivation

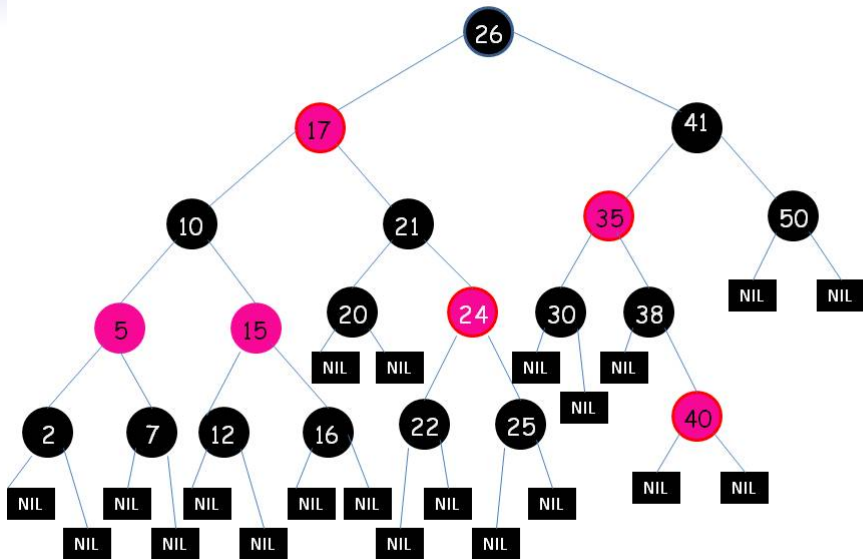
- Binary search trees support SEARCH, MINIMUM, MAXIMUM, PREDECESSOR, SUCCESSOR, INSERT and DELETE.
- Each operation takes time $O(h)$ where h is the height of the tree.
- *Notable deficit:* Binary search trees are not height balanced. So, h is not necessarily $O(\log n)$.

Summary

- *Red-Black* trees are *binary search trees* that satisfy additional properties.
- These properties imply that $h = O(\log n)$.
- Simple enough that INSERT and DELETE operations can still be done in time $O(h) = O(\log n)$.

Definition

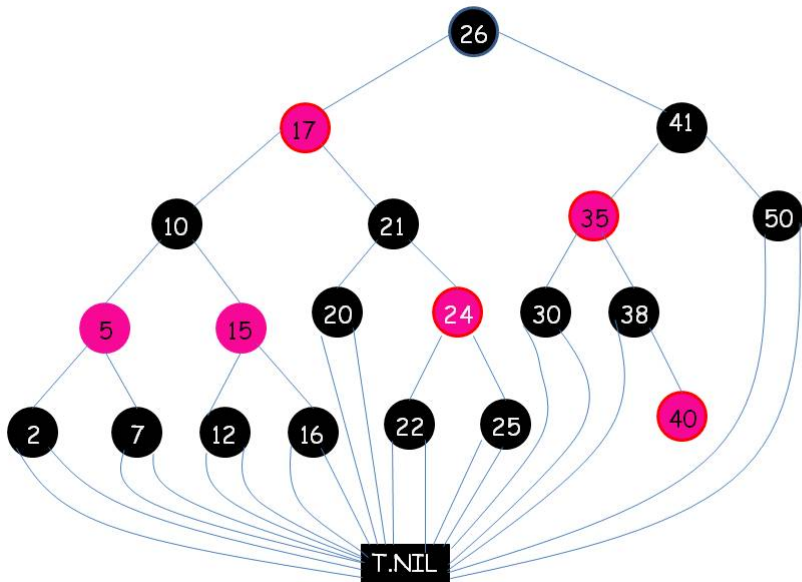
1. A red-black tree *is a **binary search tree*** with
 - one extra bit of storage per node: its color.
 - Every node is colored either **RED** or BLACK.
2. The root is black.
3. Every leaf (NIL) node is black.
4. If a node is red, then both its children are black.
5. **Black-height property:** for each node, **all** simple paths from the node to descendant leaves (NIL) contain the same number of black nodes.



An example red-black tree

Convention

- Leaf nodes are NIL nodes that are colored black (as shown earlier).
- *A slight modification:* For succinctness, a single node called $T.NIL$ is created.
 - Its color is **black**.
 - It has arbitrary values of *left*, *right* and *key*.
 - All pointers to NIL are replaced by pointers to $T.NIL$.
- This reduces storage space for the NIL nodes.



An example red-black tree with sentinel node T.NIL

Black-height of a node

- Recall Definition property 5.
 - **Black-height property:** for each node, **all** simple paths from the node to descendant leaves (NIL) contain the same number of black nodes.
- **Black-height** of a node is the number of black nodes in any simple path from, but not including, a node to a leaf node (NIL or $T.NIL$).
- The notion of *black-height* of a node is well defined by property 5, since,
 - all descending simple paths from the node have the same number of black nodes.

Black-height of node and tree

- Denote black-height of a node x as $bh(x)$.
- The black height of a red-black tree is the black height of its root node.

Near-balanced property of red-black trees-I

Property 1: The subtree rooted at any node x contains at least $2^{bh(x)} - 1$ internal nodes

- Recall: Only NIL(or $T.NIL$) nodes are leaf nodes, all other nodes are internal nodes.
- Property 1 shows a certain minimum population in the sub-tree rooted at x .
- We will prove by induction on height of a node (not black-height of a node).
- Base case: $bh(x) = 0$. This means x is a NIL(or $T.NIL$) node. Then,

$$2^{bh(x)} - 1 = 2^0 - 1 = 0$$

which is correct since x is a NIL node (no internal nodes in subtree rooted at x .)

Property 1 (proof: induction case)

- Induction case. Suppose x has positive height and is an internal node.
- Without loss of generality, let x have two children.
- The black height of a child is either $bh(x)$ (if the child is red) or $bh(x) - 1$ (if child is black).
- Height of each child of $x \leq$ height of x (trivial!).

Property 1 (proof: induction case)

- So by induction hypothesis applied to each child node of x the sub-tree rooted at each child has at least $2^{bh(x)-1} - 1$ internal nodes.

- The sub-tree rooted at x contains at least

$$\left(2^{bh(x)-1} - 1\right) + \left(2^{bh(x)-1} - 1\right) + 1 = 2^{bh(x)} - 1$$

internal nodes.

- This proves property 1: Number of internal nodes in the subtree rooted at x is at least $2^{bh(x)} - 1$.

Property 2

- Definition property 4: Red nodes have both children black.
- Half the nodes on any simple path from root to leaf (NIL) is black.
- Let h be the height of the tree (i.e., of root node).
- Hence, $h \leq 2bh(x)$ or, $bh(x) \geq h/2$.
- Hence, number of nodes

$$n \geq 2^{bh(\text{root})} - 1 \geq 2^{h/2} - 1$$

or,

$$h \leq 2\log_2(n + 1)$$

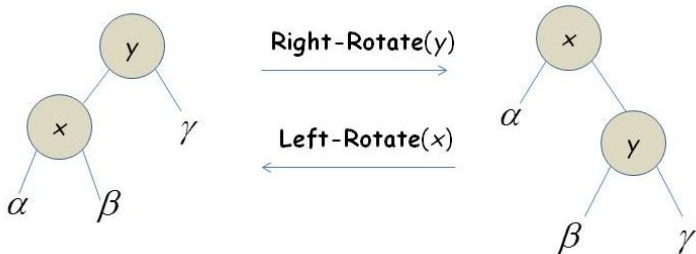
Property 2

Property 2: A red-black tree with n internal nodes has height at most $2 \log_2(n + 1)$.

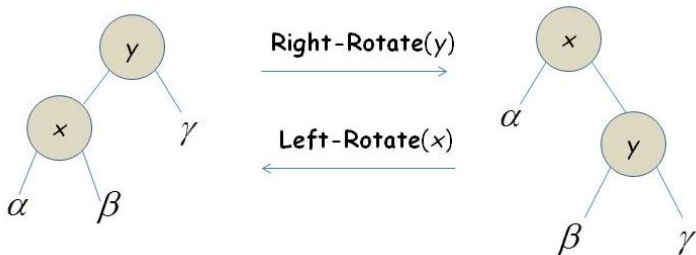
- We have just proved Property 2.
- This gives a “near-balanced” property of red-black trees.
- Since, red-black trees are binary search trees, hence, finding MAXIMUM, MINIMUM, PREDECESSOR and SUCCESSOR can be done in time $O(h) = O(\log n)$.
- We will have to design **new algorithms** for INSERT and DELETE.

Rotations: Towards Insertions and Deletions

- We will move towards designing algorithms for INSERT and DELETE operations.
- Central to this is the concept of **rotations**.
- Two kinds of rotations: Left rotation \Leftrightarrow Right rotation.

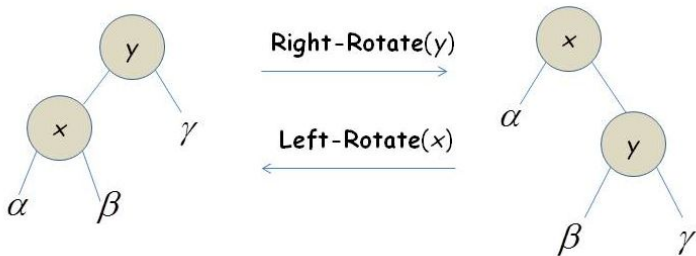


Rotations



- Figure shows right rotation of the subtree on the left about the node y and left rotation of the subtree on the right about the node x .
- α, β and γ are sub-trees (that could possibly be NIL).
- *Rotations preserve the binary search tree property.*

Rotations



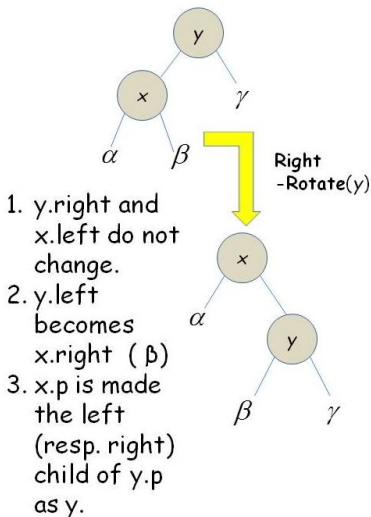
- Rotations, by themselves, do not alter the color of any node.
- After rotation, a red node may have a red child (violation!).
- Rotations may or may not preserve black height property (another possible violation!)
- But, binary search tree property is preserved!

Pseudo-code

RIGHT-ROTATE(T, y)

// Right rotate subtree rooted at y .

1. $x = y.left$
2. $y.left = x.right$
3. **if** $x.right \neq \text{NIL}$
4. $x.right.p = y$
5. $x.right = y$
6. $pp = y.p$
7. **if** $pp \neq \text{NIL}$
8. **if** $y == pp.left$
9. $x = pp.left$
10. **else**
11. $x = pp.right$
12. **else**
13. $T.root = x$
14. $x.p = pp$
15. $y.p = x$



Rotation

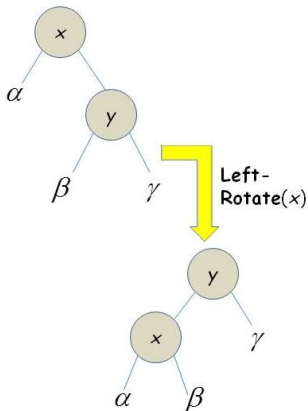
- Code for *Left-Rotate*(T, x) is symmetric.
- Both procedures run in time $O(1)$.

Left rotation: Psuedo-code

LEFT-ROTATE(T, x)

// Left rotate subtree rooted at x .

1. $y = x.right$
2. $x.right = y.left$
3. **if** $y.left \neq \text{NIL}$
4. $x.right.p = x$
5. $y.left = x$
6. $pp = x.p$
7. **if** $pp \neq \text{NIL}$
8. **if** $x == pp.left$
9. $y = pp.left$
10. **else**
11. $y = pp.right$
12. **else**
13. $T.root = y$
14. $y.p = pp$
15. $x.p = y$



Red-black tree: Insertion

- $\text{INSERT}(T, z)$.
- Two steps: one pass down the tree and another pass (possibly) back up the tree.
- In the first step, the new node, say z , is inserted as if the tree was simply a binary search tree and disregarding the coloring properties.
- The newly inserted node is colored **RED**.
- May be possible violations of the coloring properties.
- A new routine RB-INSERT-FIXUP is called to fix the problem and remove the violations.

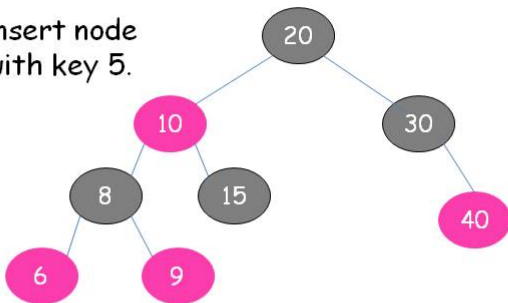
Outline of insertion

RB-INSERT(T, z)

1. BSTINSERT(T, z)
2. $z.color = \text{RED}$
3. RB-INSERT-FIXUP(T, z)

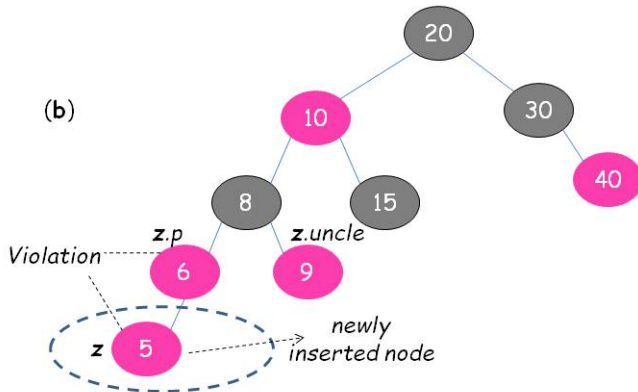
Example

Insert node
with key 5.



Insert as if the tree was a binary search tree, and insert the node as a leaf with red color.

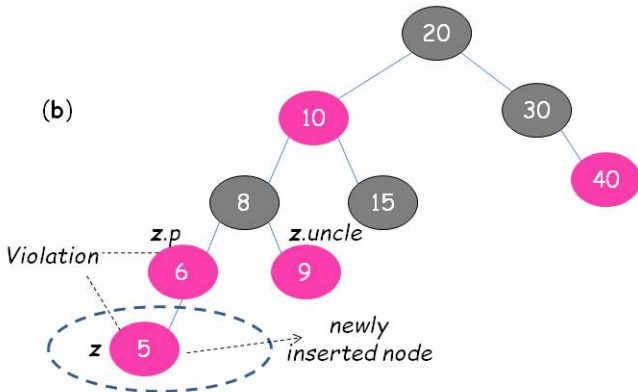
Example contd.



After insertion as in a binary search tree

- The first step is to insert the node z as if the red-black tree is a binary search tree.
- That is, insert into the position just “as z falls off the tree”.
- Color z as red.
- *The black-height of all nodes remain unchanged.*
- If $z.p$ (z 's parent) is black, insertion is completed.
- If $z.p$ is red, then both z and its parent are red: this is a violation. So $\text{RB-INSERT-FIXUP}(T, z)$.

Example contd.



- Case 1 applies: z is red, z 's parent is red and z 's uncle is red.
- Make z 's parent and z 's uncle black and make z 's grandparent red. Set z to z 's grandparent.

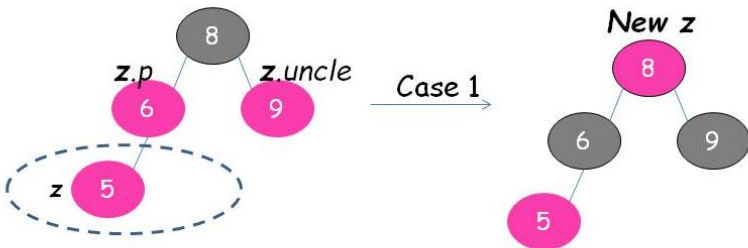
Case 1

RB-Insertion: Case 1 out of three cases.

- z is colored red, $z.p$ is colored red — Basic violation, and,
- $z.uncle$ is colored red.
- Therefore, $z.p.p$ (grandparent of z) exists.
- Why? Because, $z.p$ is colored red and cannot be the root of T . Root is black.
- $z.p.p$ is colored black.
- Why? It has two red children $z.p$ and $z.uncle$ and so $z.p.p$ cannot be red (red nodes must have only black children). So it is black.

Case 1

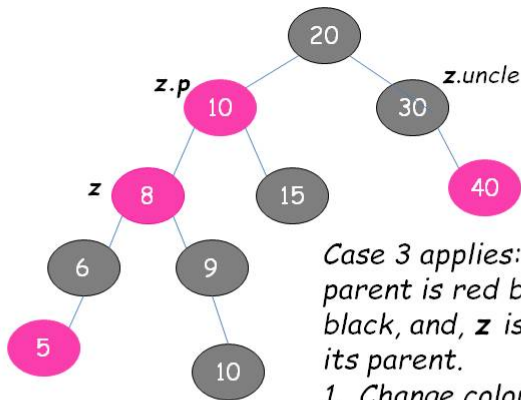
- z is colored red, $z.p$ is colored red and $z.uncle$ is colored red. $z.p.p$ is black.
- Change color of $z.p$ and $z.uncle$ to black. Change color of $z.p.p$ to red.



Pushes any possible violation up by two levels

Case 3

(c)



*Case 3 applies: **z** is red, **z**'s parent is red but **z**'s uncle is black, and, **z** is the left child of its parent.*

- 1. Change color of **z.p** to black.*
- 2. Change color of **z**'s grandparent to red.*
- 3. Rotate **right** about **z.p.p**.*

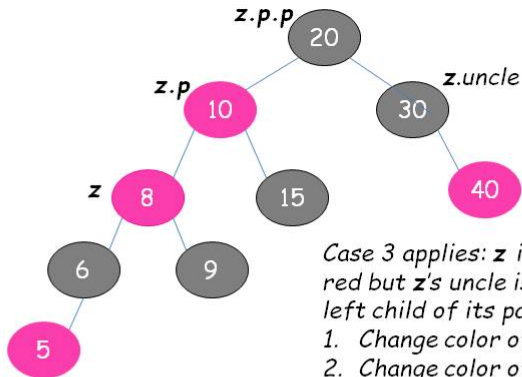
Case 3: definition

- z is colored red, and,
- $z.p$ is colored red, and,
- $z.uncle$ is colored black, and,
- z is left child of its parent.

Case 3: Handling

- Change color of $z.p$ to black.
- Change color of z 's grandparent to red.
- Rotate right about $z.p.p$.

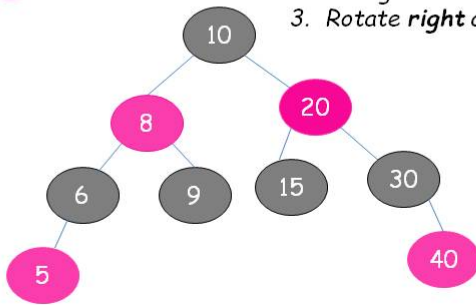
(c)



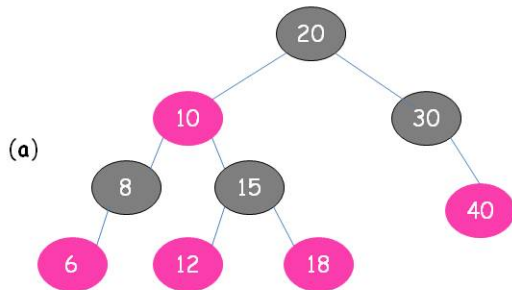
Case 3 applies: **z** is red, **z**'s parent is red but **z**'s uncle is black, and, **z** is the left child of its parent.

1. Change color of **z.p** to black.
2. Change color of **z.p.p** to **red**.
3. Rotate **right** about **z.p.p**.

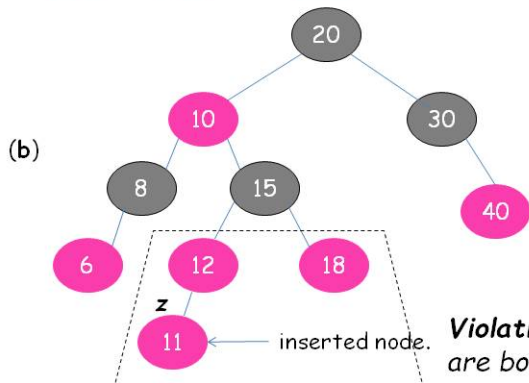
(d)



This is a proper red-black tree with no violations.



Insert node with key 11.



Case 1 applies: *z's uncle is red.*

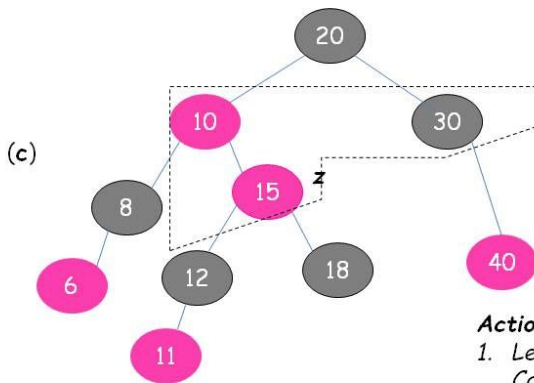
1. **Make z's parent and z uncle black.**

2. **Make z's grandparent red.**

3. **Set z to be z's grandparent.**

Violation: *z and z's parent are both colored red.*

Case 2



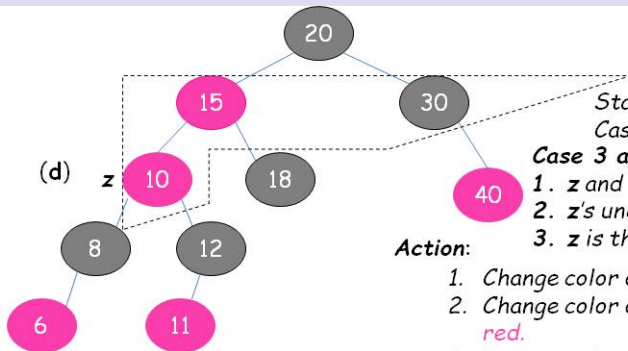
State of tree after Case 1 is processed.
Current focus shown in dots.

Case 2 applies because:

1. **z** and **z**'s parent are red.
2. **z**'s uncle is black.
3. **z** is the right child of its parent.

Action:

1. Left-rotate about **z** to get to Case 3.
2. **Set z** to its current parent.



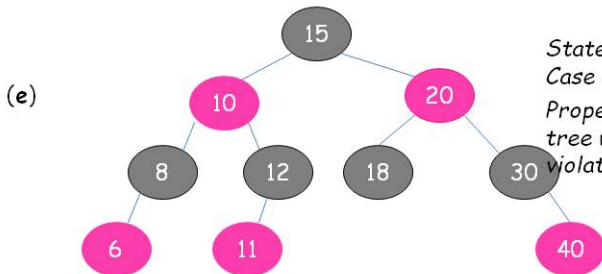
State of tree after
Case 2 is processed.

Case 3 applies because:

1. *z* and *z*'s parent are red.
2. *z*'s uncle is black.
3. *z* is the left child of its parent.

Action:

1. Change color of *z.p* to black.
2. Change color of *z*'s grandparent to *red*.
3. Rotate **right** about *z.p*.



State of tree after
Case 3 is processed.

Proper red-black
tree with no
violations.

Insert Fixup routine

- Repeatedly, check which case the violation falls into, and apply the corresponding fixup.

RB-INSERT-FIXUP(T, z)

1. $caseno = \text{RB-INSERT-FIXUP-CHECK-CASE}(T, z)$
2. **while** $caseno > 0$
3. **if** $caseno == 1$
4. $\text{RB-INSERT-FIXUP-APPLY-CASE-1}(T, z)$
5. **elseif** $caseno == 2$
6. $\text{RB-INSERT-FIXUP-APPLY-CASE-2}(T, z)$
7. **elseif** $caseno == 3$
8. $\text{RB-INSERT-FIXUP-APPLY-CASE-3}(T, z)$
9. $caseno = \text{RB-INSERT-FIXUP-CHECK-CASE}(T, z)$
10. $T.root.color = \text{BLACK}$

Some observations

- Fixup of Case 2 leads to Case 3.
- Fixup of Case 3 gives a red-black tree with no violations: hence insertion terminates.
- Fixup of Case 1 pushes the violation up the tree by two levels. This can be iterative.
- Fixups of all cases are done in $O(1)$ time.
- Number of times Case 1 repeatedly occurs is $O(h) = O(\log n)$. Hence, time taken is $O(\log n)$.

Auxiliary procedure

UNCLE(z)

1. **if** $z == \text{NIL}$ **or** $z.p == \text{NIL}$ **or** $z.p.p == \text{NIL}$
2. **return** NIL
3. **elseif** $z.p == z.p.p.\text{left}$
4. **return** $z.p.p.\text{right}$
5. **else**
6. **return** $z.p.p.\text{left}$

Find which case is violated

RB-INSERT-FIXUP-CHECK-CASE(T, z)

7. **if** $z.color == \text{BLACK}$ **or** $z.p == \text{NIL}$ **or** $z.p.color == \text{BLACK}$
8. **return** 0
9. **elseif** $\text{UNCLE}(z).color == \text{RED}$
10. **return** 1
11. **elseif** $z == z.p.right$
12. **return** 2
13. **else**
14. **return** 3

Pseudo-code: RB-INSERT-FIXUP

RB-INSERT-FIXUP(T, z)

1. $caseno = \text{RB-INSERT-FIXUP-CHECK-CASE}(T, z)$
2. **while** $caseno > 0$
3. **if** $caseno == 1$
4. $z.p.color = \text{BLACK}$ // Case 1
5. $\text{UNCLE}(z).color = \text{BLACK}$ // Case 1
6. $z.p.p.color = \text{RED}$ // Case 1
7. $z = z.p.p$ // Case 1
8. **elseif** $caseno == 2$
9. $\text{LEFT-ROTATE}(T, z)$ // Case 2
10. $z = z.left$ // Case 2
11. **elseif** $caseno == 3$
12. $z.p = \text{BLACK}$ // Case 3
13. $z.p.p = \text{RED}$ // Case 3
14. $\text{RIGHT-ROTATE}(T, z.p)$ // Case 3
15. $T.root.color = \text{BLACK}$

Invariant

Let z denote the node that is inserted as per insertion into binary search tree. z is colored red. Invariant:

- a. z is colored red.
- b If $z.p$ is the root, then $z.p$ is black.
- c If the tree violates any of the red-black properties then it violates at most one, namely, either
 - c.1 Property 2: The root is black, or,
 - c.2 Property 4: If a node is red, then both its children are black.

Initialization of the RB-INSERT-FIXUP-PROCEDURE

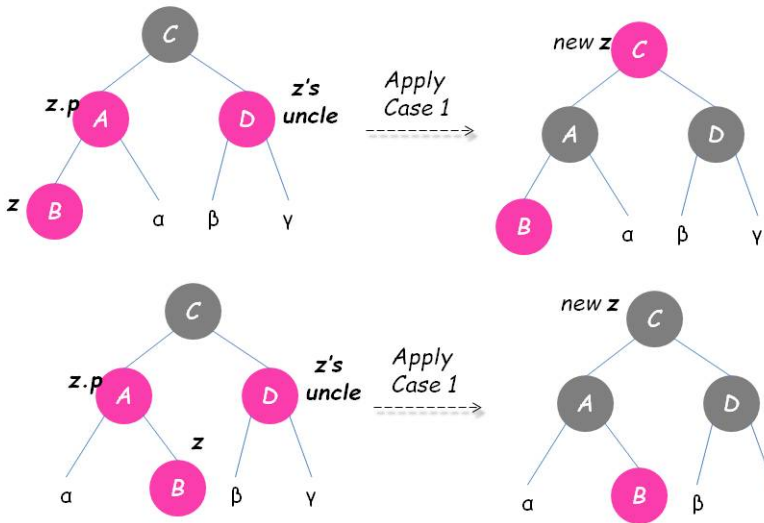
- When this procedure is initialized, z is the inserted node and is colored red.
- If $z.p$ is the root, then $z.p$ started out black and did not change prior to call.
- Important: black-heights of all old nodes remain unchanged, because z is inserted as a red node.
- If root is red, then, the root must be z . (This case is easy: change z 's color to black and terminate.)
- Otherwise, the only possible violation is that $z.p$ and z are both red. No other red-black properties are violated anywhere else in the tree.

Invariant: Maintenance

- Consider the more complicated case when z and $z.p$ are colored red, and therefore, $z.p.p$ exists and is colored black.
- There are two symmetric cases:
 1. $z.p$ is the left child of $z.p.p$ (considered here), and,
 2. $z.p$ is the right child of $z.p.p$ (symmetric, not considered).

Invariant: Maintenance

- **Case 1:** z 's uncle y is red.
- So z , $z.p$ and $z.uncle$ are all red.
- Action:
 1. Change color of $z.p$ and $z.uncle$ to black.
 2. Change color of $z.p.p$ to red.
 3. Set z to $z.p.p$. Make the grandparent as the new z .



Case 1: z is red, z 's parent is red and z 's uncle is red. The action taken is to make the parent and the uncle of z to be black and make the grandfather of z to be red. This may "**push the violation up the tree.**" There is a symmetric case when $z.p$ is the right child of its parent.

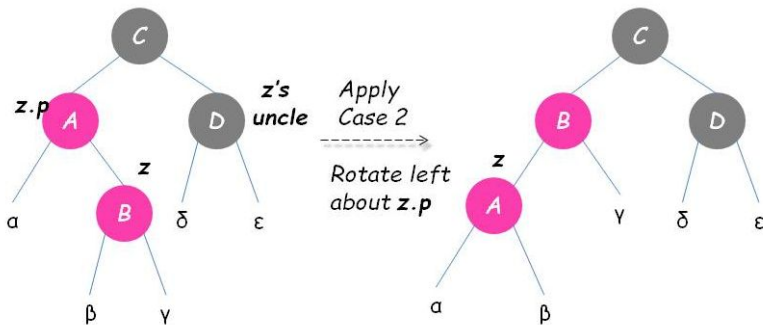
Invariant: Maintenance

- **Note:** Black height of every node remains the same before and after the transformation corresponding to Case 1.
- There can be at most one violation, namely that the new z (which is the old grandparent) is red and *its* parent may also be colored red.
- In this case, the violation moves up the tree by two levels.
- There are no other possible violations in the red-black tree.
- If the new z 's parent is colored black, then the process terminates.

Invariant: Maintenance

- **Case 2:** z' uncle y is black and z is the right child of its parent.
- **Case 3:** z 's uncle y is black and z is the left child of its parent.

Case 2



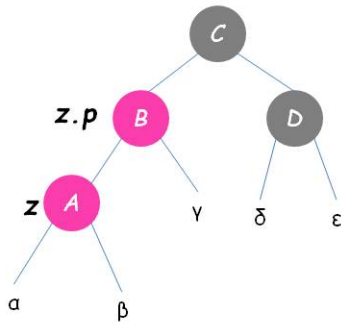
Case 2: There is a symmetric case for Case 2 when z 's parent is the right child of its parent.

Case 2

- Action: Rotate left around $z.p$. z becomes the parent and $z.p$ becomes the left child of z .
- Rename $z.p$ as z .
- We are now in Case 3. (The purpose of the action for case 2 is to transform into an instance of Case 3).

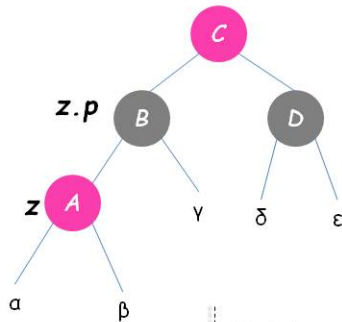
Case 3

- z and $z.p$ are red, z 's uncle is black and z is the left child of $z.p$.
- Note that $z.p.p$ is black.
- Actions:
 1. Change the color of $z.p$ and $z.uncle$ to black.
 2. Change color of $z.p.p$ to red.
 3. Rotate right about $z.p.p$.



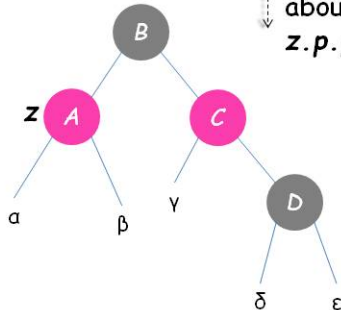
Apply
Case 3

Change
colors of B, D
and C



1. The black heights of A, α , β , γ and D are same.
2. So black height property of A, C, B are preserved.
3. Also, the black height of B is the same as the black height of C in the original tree.

Rotate
about
z.p.p



Invariant: Case 3

- Case 3 preserves black height property (needs some verification).
- There are no violations of red-black properties.
- So, the insertion procedure terminates.
- Case 3 is a termination case.

Complexity Analysis

- Each of the cases takes $O(1)$ time to complete.
- Case 1 may be called at most $O(\log n)$ times.
- If Case 2 or Case 3 is called, then the procedure terminates in $O(1)$ time.
- Hence time required is bounded by $O(\log n)$.